SDS 4030 - Statistics for Data Science II: Final Project

Main Markdown File

Noah Calabrese, Natalia Jamula, Tim Mellman

Preliminary library loading, reading code in, and exploratory data analysis

```r
gss = read.table("gss.txt", header = TRUE)
summary(gss)
```

```
##       Age             Educ           Hours           Sex
##  Min.   :18.00   Min.   : 4.00   Min.   : 3.00   Length:815
##  1st Qu.:32.00   1st Qu.:12.00   1st Qu.:39.00   Class :character
##  Median :42.00   Median :14.00   Median :40.00   Mode  :character
##  Mean   :42.91   Mean   :13.93   Mean   :42.59
##  3rd Qu.:52.00   3rd Qu.:16.00   3rd Qu.:50.00
##  Max.   :88.00   Max.   :20.00   Max.   :89.00
##     TVHours          Happy               MEduc           FEduc
##  Min.   : 0.000   Length:815         Min.   : 0.00   Min.   : 0.00
##  1st Qu.: 2.000   Class :character   1st Qu.:10.00   1st Qu.: 9.00
##  Median : 2.000   Mode  :character   Median :12.00   Median :12.00
##  Mean   : 2.709                      Mean   :11.61   Mean   :11.45
##  3rd Qu.: 3.000                      3rd Qu.:14.00   3rd Qu.:14.00
##  Max.   :20.000                      Max.   :20.00   Max.   :20.00
##    Marital               Sibs            Work              Zodiac
##  Length:815         Min.   : 0.000   Length:815         Length:815
##  Class :character   1st Qu.: 2.000   Class :character   Class :character
##  Mode  :character   Median : 3.000   Mode  :character   Mode  :character
##                     Mean   : 3.439
##                     3rd Qu.: 5.000
##                     Max.   :16.000
##      Income
##  Min.   :     500
##  1st Qu.:   23750
##  Median :   45000
##  Mean   :  324767
##  3rd Qu.:   67500
##  Max.   :28637626
```

We begin our data science workflow by familiarizing ourselves with the GSS dataset. The dataset contains 815 rows corresponding to 815 individuals' survery responses. The dataset contains 13 variables. There are 8 discrete numerical variables: Age (Responder's Age), Educ (Total years of education), Hours (Weekly work hours), TVHours (Hours spent watching TV daily), MEduc (Mother's

years of education), FEduc (Father's years of education), Sibs (Number of siblings), and Income (Yearly income). Thus there are 5 categorical variables: Sex (Responder's sex), Happy (Happiness level), Marital (Marital status), Work (Work status), and Zodiac (Zodiac sign). Uniquely they take on the following values: Sex (Female, Male), Happy (Pretty Happy, Very Happy, Not Too Happy), Marital ("Divorced, Widowed, Married, Never Married, Separated), Work (Fulltime, Housekeeping, Laid Off, School, Retired, Parttime, Not Working, Other), and Zodiac (Virgo, Libra, Taurus, Pisces, Cancer, Aquarius, Leo, Capricorn, Aries, Gemini, Sagittarius, Scorpio).

Now that we understand our data we isolate our goal being to build a linear model for the response variable yearly income by treating all other seven quanitative variables as covariates and ignoring all five qualitative variables.

Step 1:

i) Fit a multiple linear model with all available quantitative covariates.

ii) Perform exact and asymptotic tests for the set of linear constraints which has been specified for your team.

iii) Do all of your conclusions agree with each other and what are the relationships between the different tests?

---

i) Fitting a multiple linear model with all available quantitative covariates. This is the model that we will use for the next few parts of this project

```r
library(faraway)
full_model <- lm(Income ~ Age + Educ + Hours + TVHours + MEduc + FEduc +
    Sibs, gss)
summary(full_model)
```

```
##
## Call:
## lm(formula = Income ~ Age + Educ + Hours + TVHours + MEduc +
##     FEduc + Sibs, data = gss)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1211280  -410167  -271382  -129856 28149890
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)   211943     688636   0.308   0.7583
## Age             2949       6231   0.473   0.6361
```

```
## Educ           20448      35626   0.574   0.5661
## Hours          -9324       5923  -1.574   0.1159
## TVHours       -65363      42493  -1.538   0.1244
## MEduc          -3715      32699  -0.114   0.9096
## FEduc          10206      28724   0.355   0.7225
## Sibs           58649      33490   1.751   0.0803 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2340000 on 807 degrees of freedom
## Multiple R-squared:  0.0101, Adjusted R-squared:  0.001516
## F-statistic: 1.177 on 7 and 807 DF,  p-value: 0.3136
```

For the sake of later model comparison we will find the least squares estimate of the coefficient vector, an unbiased estimate of the residual variance and the MLE of the residual variance in the full model.

```r
Y = gss$Income
X = model.matrix(full_model)
n = dim(X)[1]
p = dim(X)[2]
beta = solve(crossprod(X), crossprod(X, Y))[, ]
print(beta)
```

```
## (Intercept)         Age        Educ       Hours     TVHours       MEduc
##  211942.841    2949.415   20448.463   -9323.480  -65363.112   -3715.210
##       FEduc        Sibs
##   10205.481   58648.969
```

```r
residuals = Y - X %*% beta
sigma = sqrt(mean(residuals^2))
print(sigma^2)
```

```
## [1] 5.42076e+12
```

```r
S = sqrt(sum(residuals^2)/(n - p))
print(S^2)
```

```
## [1] 5.474497e+12
```

Our full model is $Y_i = \beta_0 + \beta_{Age}\text{Age}_i + \beta_{Educ}\text{Educ}_i + +\beta_{Hours}\text{Hours}_i + \beta_{TVHours}\text{TVHours}_i + \beta_{MEduc}\text{MEduc}_i + \beta_{FEduc}\text{FEduc}_i + \beta_{Sibs}\text{Sibs}_i + \epsilon_i$. An unbiased estimate of the residual variance is equal to 5.42076e+12 and the MLE of the residual variance is equal to 5.474497e+12, under the full model.

The reduced model is a little more tricky where we have to create a model in which it follows the constraints given by: $H_0 : 10 * \beta_{MEduc} = 20 * \beta_{FEduc} = 2$.

We can rewrite the constraint as $10 * \beta_{MEduc} = 2$, $20 * \beta_{FEduc} = 2$. Or, $\beta_{MEduc} = 0.2$, $\beta_{FEduc} = 0.1$

We can plug the null values into the full model, replacing coefficients with the constants so the full model becomes $Y_i = \beta_0 + \beta_{Age}\text{Age}_i + \beta_{Educ}\text{Educ}_i + \beta_{Hours}\text{Hours}_i + \beta_{TVHours}\text{TVHours}_i + 0.2 * \text{MEduc}_i + 0.1 * \text{FEduc}_i + \beta_{Sibs}\text{Sibs}_i + \epsilon_i$

We have to set MEduc and FEduc, the fixed parts of the response in the reduced model, as the offset.

```
reduced_model <- lm(Income ~ Age + Educ + Hours + TVHours + Sibs, gss,
    offset = 0.2 * MEduc + 0.1 * FEduc)
summary(reduced_model)
```

```
##
## Call:
## lm(formula = Income ~ Age + Educ + Hours + TVHours + Sibs, data = gss,
##     offset = 0.2 * MEduc + 0.1 * FEduc)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1208997  -397497  -267841  -139546 28119340
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)   286038     616017   0.464   0.6425
## Age             2334       5805   0.402   0.6878
## Educ           23374      33141   0.705   0.4808
## Hours          -9411       5911  -1.592   0.1117
## TVHours       -66218      42346  -1.564   0.1183
## Sibs           56138      32385   1.733   0.0834 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2337000 on 809 degrees of freedom
## Multiple R-squared:  0.009939,   Adjusted R-squared:  0.00382
## F-statistic: 1.624 on 5 and 809 DF,  p-value: 0.151
```

```
R <- rbind(c(0, 0, 0, 0, 0, 1, 0, 0), c(0, 0, 0, 0, 0, 0, 1, 0))  ##the 1's in R represent the
r <- c(0.2, 0.1)
k <- length(r)
beta0 = beta - solve(crossprod(X), crossprod(R, solve(R %*% solve(crossprod(X),
    t(R)), R %*% beta - r)))[, ]
print(beta0)
```

```
## (Intercept)         Age        Educ       Hours     TVHours       MEduc
```

```
##  286037.934    2333.632   23374.285   -9411.360  -66218.159       0.200
##       FEduc        Sibs
##       0.100   56138.450
```

```
residuals0 = Y - X %*% beta0
sigma0 = sqrt(mean(residuals0^2))
print(sigma0^2)
```

```
## [1] 5.421657e+12
```

```
S0 = sqrt(sum(residuals0^2)/(n - p + k))
print(S0^2)
```

```
## [1] 5.461868e+12
```

An unbiased estimate of the residual variance is equal to 5.421657e+12 and the MLE of the residual variance is equal to 5.461868e+12, under the reduced model.

ii) We have to test the hypothesis $H_0 : 10 * \beta_{MEduc} = 20 * \beta_{FEduc} = 2$ against exact and asymptotic tests. The best exact test is, of course, the exact F-test assuming we are have normalcy conditions (which we will check later, but step 1 ensures that we do it regardless of the normalcy checks). To test for the asymptotic tests it is logical to use the Wald test since we have a large sample (n = 815). It is also one of the easier asymptotic tests to check, so we will use it due to its efficiency, accuracy, and sample size.

Exact F- Test

```
FT = crossprod(R %*% beta - r, solve(R %*% solve(crossprod(X), t(R)),
    R %*% beta - r))[, ]/(k * S^2)
print(FT)
```

```
## [1] 0.06680263
```

```
pvalue = pf(FT, k, n - p, lower.tail = FALSE)
print(pvalue)
```

```
## [1] 0.935385
```

```
anova(reduced_model, full_model)
```

```
## Analysis of Variance Table
##
## Model 1: Income ~ Age + Educ + Hours + TVHours + Sibs
## Model 2: Income ~ Age + Educ + Hours + TVHours + MEduc + FEduc + Sibs
##   Res.Df        RSS Df  Sum of Sq      F Pr(>F)
## 1    809 4.4187e+15
## 2    807 4.4179e+15  2 7.3142e+11 0.0668 0.9354
```

The observed value of the F test statistic is equal to 0.06680263 and the corresponding p-value is equal to 0.935385, so we fail to reject the null hypothesis at significance level $\alpha = 5\%$.

Wald Test

```r
WT = crossprod(R %*% beta - r, solve(R %*% solve(crossprod(X), t(R)),
    R %*% beta - r))[, ]/S^2
print(WT)
```

```
## [1] 0.1336053
```

```r
pvalue = pchisq(WT, k, lower.tail = FALSE)
print(pvalue)
```

```
## [1] 0.9353798
```

```r
anova(reduced_model, full_model, test = "Chisq")
```

```
## Analysis of Variance Table
##
## Model 1: Income ~ Age + Educ + Hours + TVHours + Sibs
## Model 2: Income ~ Age + Educ + Hours + TVHours + MEduc + FEduc + Sibs
##   Res.Df        RSS Df  Sum of Sq Pr(>Chi)
## 1    809 4.4187e+15
## 2    807 4.4179e+15  2 7.3142e+11   0.9354
```

The observed value of the Wald test statistic is equal to 0.1336053 and the corresponding p-value is equal to 0.9353798, so we fail to reject the null hypothesis at significance level $\alpha = 5\%$

Score Test

```r
SF = crossprod(X, residuals0)/sigma0^2
FI = crossprod(X)/sigma0^2
ST = crossprod(SF, solve(FI, SF))[, ]
print(ST)
```

```
## [1] 0.1349074
```

```r
P = X %*% solve(crossprod(X), t(X))
ST = sum((P %*% residuals0)^2)/sigma0^2
print(ST)
```

```
## [1] 0.1349074
```

```r
pvalue = pchisq(ST, k, lower.tail = FALSE)
print(pvalue)
```

```
## [1] 0.934771
```

The observed value of the Score test statistic is equal to 0.1349074 and the corresponding p-value is equal to 0.934771, so we fail to reject the null hypothesis at significance level $\alpha = 5\%$.

Likelihood Ratio Test

```r
LR = -2 * (logLik(reduced_model)[1] - logLik(full_model)[1])
print(LR)
```

```
## [1] 0.1349186
```

```r
LR = n * log(sigma0^2/sigma^2)
print(LR)
```

```
## [1] 0.1349186
```

```r
pchisq(LR, k, lower.tail = FALSE)
```

```
## [1] 0.9347658
```

```r
library(lmtest)
lrtest(reduced_model, full_model)
```

```
## Likelihood ratio test
##
## Model 1: Income ~ Age + Educ + Hours + TVHours + Sibs
## Model 2: Income ~ Age + Educ + Hours + TVHours + MEduc + FEduc + Sibs
##   #Df LogLik Df  Chisq Pr(>Chisq)
## 1   7 -13105
## 2   9 -13105  2 0.1349     0.9348
```

The observed value of the Likelihood Ratio test statistic is equal to 0.1349186 and the corresponding p-value is equal to 0.9347658, so we fail to reject the null hypothesis at significance level $\alpha = 5\%$.

iii) Yes, all of my conclusion agree with each other as in the exact F test, Wald test, score test, and likelihood ratio test we have failed to reject the null, and they all yield have similar p-values. Both the F-Exact test and asymptotic tests provide very small test statistics which means we do not have substantial evidence to reject the null hypothesis.

```
WT = k * FT
print(WT)
```

```
## [1] 0.1336053
```

```
ST = WT * S^2/sigma0^2
print(ST)
```

```
## [1] 0.1349074
```

```
LR = n * log(n/(n - p) * WT/ST)
print(LR)
```

```
## [1] 0.1349186
```

```
LR = n * log(1 + WT/(n - p))
print(LR)
```

```
## [1] 0.1349186
```

```
LR = -n * log(1 - ST/n)
print(LR)
```

```
## [1] 0.1349186
```

In the F test, the exact F-statistic is 0.06680263, and the p-value is 0.935385. In the Wald test the Wald statistic is 0.1336053 and the p-value is 0.9353798. According to Slutsky's theorem, we infer that W_n = kF, the Wald statistic equals kF asymptotically. The Wald statistic is k times the 0.06680263 where k=2. While the F test relies on a finite sample, the Wald statistic follows a chi-square distribution and asymptotic approximations.

In the score test, the score test statistic is 0.1349074 and the p-value is 0.934771 We know that $\hat{\sigma}_0^2 ST = S^2 W_n$. The Wald and Score test statistics only differ in terms of the residual variance estimation. The Wald test statistic utilizes an unbiased estimator of the residual variance under the full model while the Score test statistic utilizes the MLE of the residual variance under the

reduced model. In other words the Wald test uses the full model variance and the Score test uses reduced model variance. Thus we can can see how the two test statistics differ slightly.

In the likelihood ratio test, likelihood test statistic is 0.1349186 and the p-value is 0.9347658. Below we see how the Likelihood test statistic relates to the Wald test and Score test through transformations, implying they provide the same inference for very large sample sizes.

So yes, we observe that the results of all the four test statistics concur. The results of the asymptotic Wald, Likelihood Ratio and Score tests are perfectly in line with the results of the exact F test, which is expected since n = 815 which is large enough to be accurate for an asymptotic test. All three asymptotic tests, Wald, Score & Likelihood Ratio Tests, are all asymptotically equivalent, converging to the same test statistics and giving us the same conclusions. They measure deviation from the null differently, where the Wald test measures the horizontal differences of the parameters and the null, the Score test measures the slope of the log likelihood at the null, and the likelihood ratio test measures the vertical differences in log likelihood.

Step 2:

i) Is the variance of the random error terms non-constant as a function of any of the quantitative or qualitative covariates?

ii) Perform suitable heteroscedasticity tests and state your conclusions.

---

i) We begin by plotting and examining a residuals vs fitted plot of the full model (shown in part 1). We can note that the residuals spread further away from each other as the fitted values increases, creating a fan-like pattern for values above the red line, as serving as a indicator of heteroscedasticity. Additionally, it also has some extreme outliers and leverage points (also called influential outliers). The points are basically pulling the regression line and likely causing that massive variance, or at least factoring into the magnitude of it.

```r
library(faraway)
full_model <- lm(Income ~ Age + Educ + Hours + TVHours + MEduc + FEduc +
    Sibs, data = gss)
Y = gss$Income
X = model.matrix(full_model)
n = dim(X)[1]
p = dim(X)[2] - 1
residuals = full_model$residuals
plot(fitted(full_model), residuals, xlab = "Fitted Values", ylab = "Residuals",
    main = "Residuals vs Fitted")
abline(h = 0, col = "red")
```

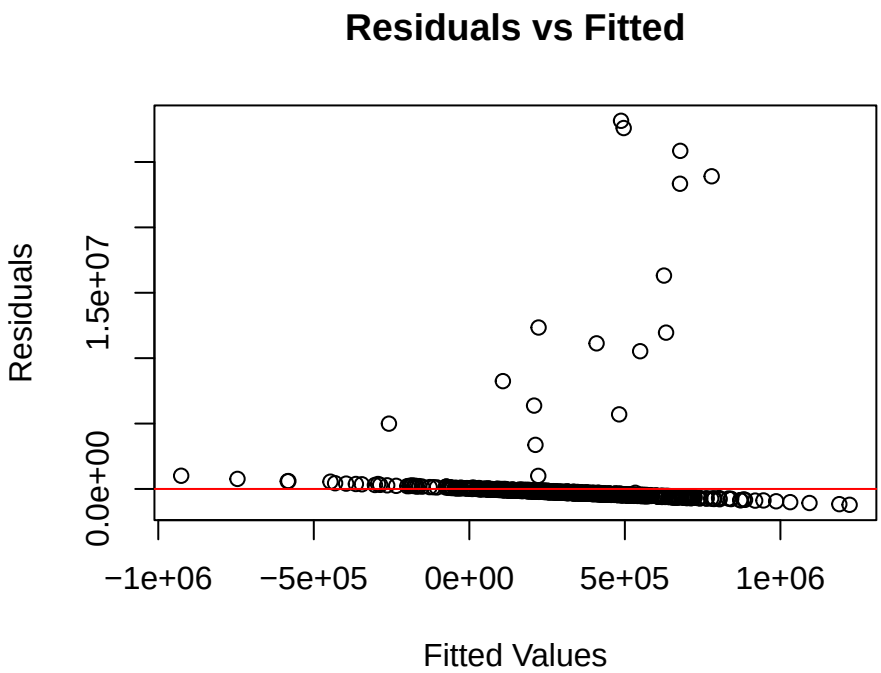

Quantitative Covariates

Responder's age

```
plot(rstandard(full_model) ~ Age, data = gss, ylab = "Standardized Residuals",
    xlab = "Age", pch = 16)
abline(h = 0, lty = 2, lwd = 2)
```

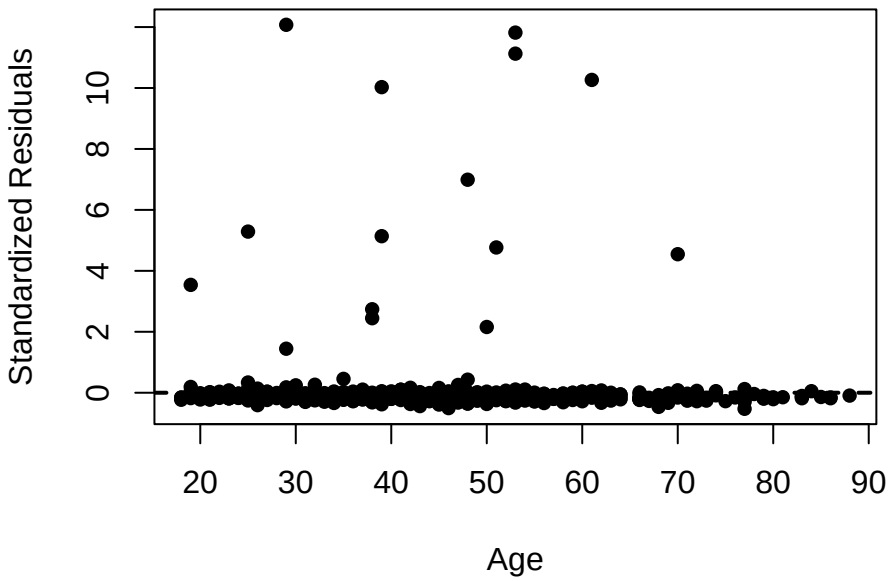

Total years of education

```
plot(rstandard(full_model) ~ Educ, data = gss, ylab = "Standardized Residuals",
    xlab = "Educ", pch = 16)
abline(h = 0, lty = 2, lwd = 2)
```

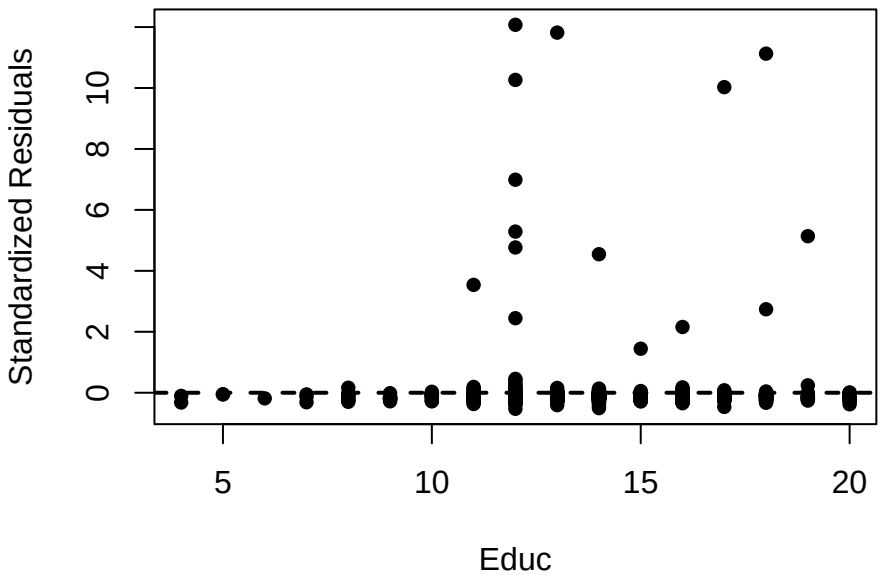

Weekly work hours

```
plot(rstandard(full_model) ~ Hours, data = gss, ylab = "Standardized Residuals",
    xlab = "Hours", pch = 16)
abline(h = 0, lty = 2, lwd = 2)
```

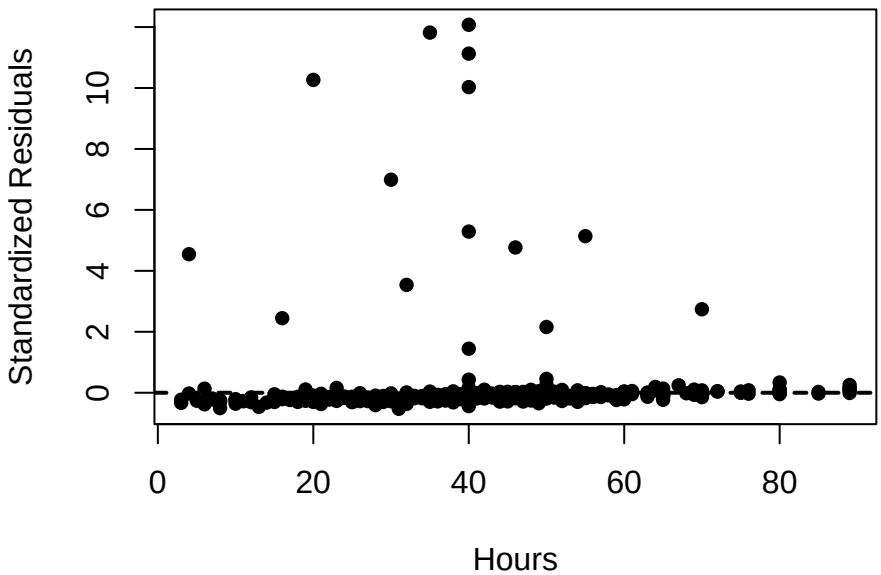

Hours spent watching TV daily

```
plot(rstandard(full_model) ~ TVHours, data = gss, ylab = "Standardized Residuals",
    xlab = "TVHours", pch = 16)
abline(h = 0, lty = 2, lwd = 2)
```

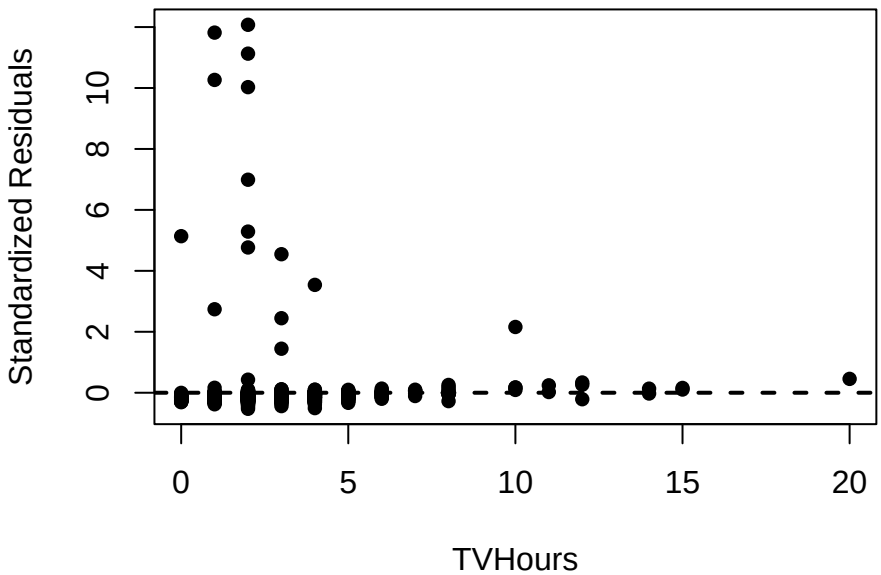

Mother's years of education

```
plot(rstandard(full_model) ~ MEduc, data = gss, ylab = "Standardized Residuals",
    xlab = "MEduc", pch = 16)
abline(h = 0, lty = 2, lwd = 2)
```

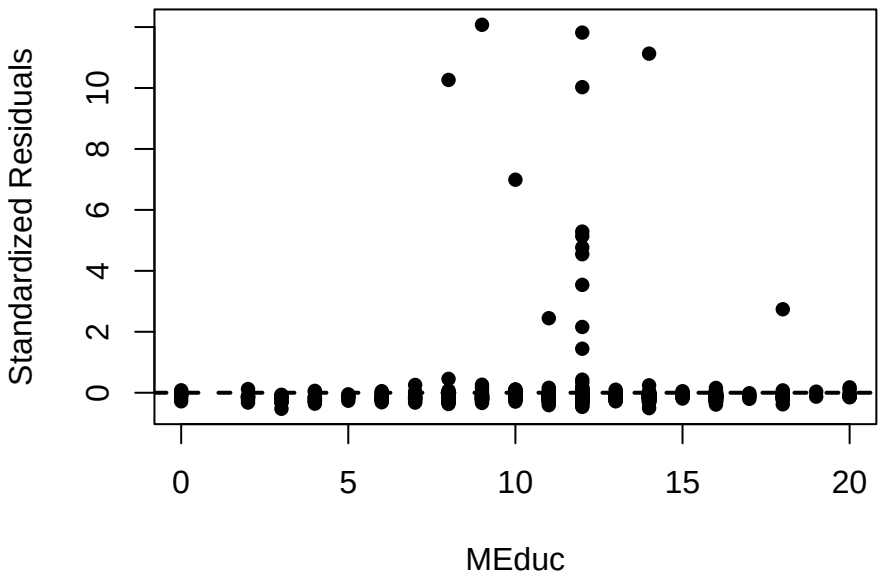

Father's years of education

```
plot(rstandard(full_model) ~ FEduc, data = gss, ylab = "Standardized Residuals",
    xlab = "FEduc", pch = 16)
abline(h = 0, lty = 2, lwd = 2)
```

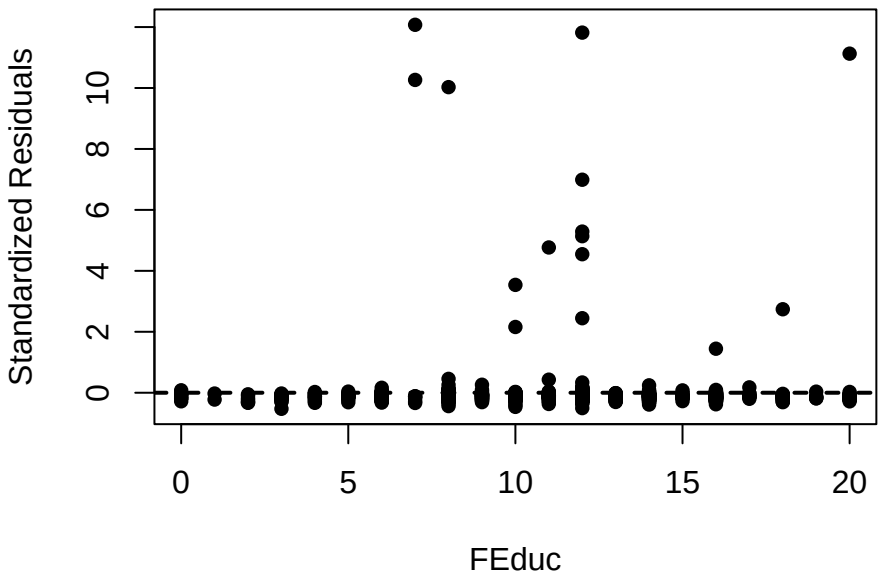

Number of siblings

```
plot(rstandard(full_model) ~ Sibs, data = gss, ylab = "Standardized Residuals",
    xlab = "Sibs", pch = 16)
abline(h = 0, lty = 2, lwd = 2)
```

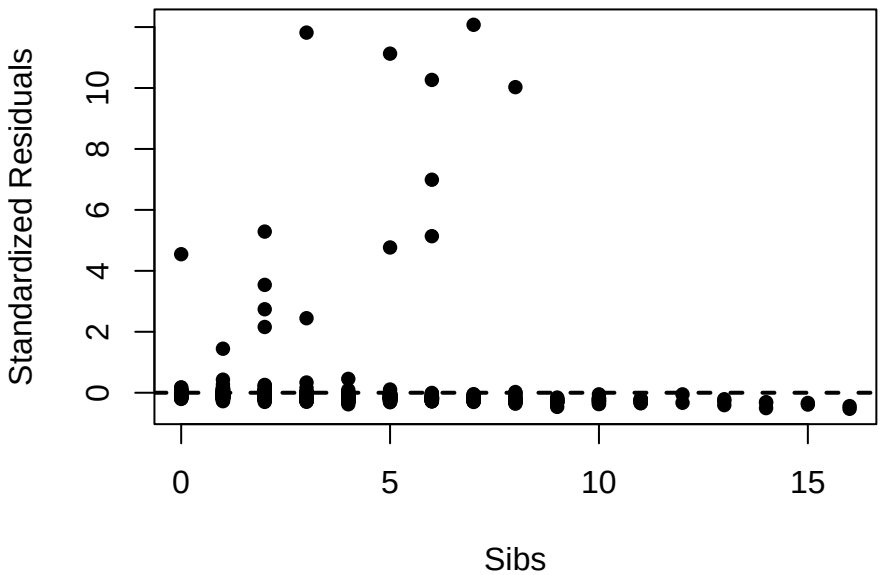

After inspecting the standardized residuals plotted against each quantitative covariate, I noticed the presence of extreme outliers but no systematic pattern of increasing or decreasing spread as a function of a quantitative covariate that would indicate the form of heteroscedasticity that appears when the variance of the random error terms is a function of some of the available covariates.

Qualitative Covariates

Responder's sex

```
boxplot(rstandard(full_model) ~ Sex, data = gss, ylab = "Standardized Residuals",
    xlab = "Sex")
abline(h = 0, lty = 2, lwd = 2)
```

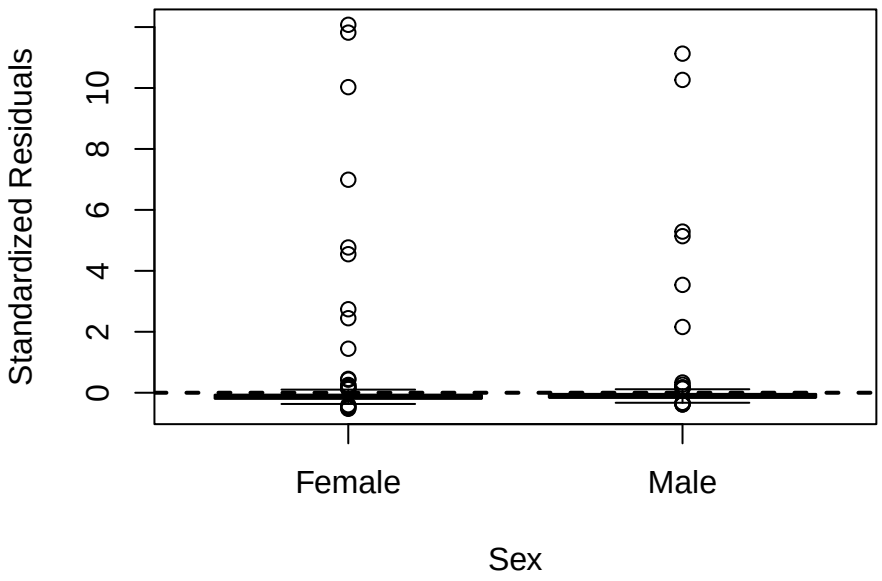

Happiness level

```
boxplot(rstandard(full_model) ~ Happy, data = gss, ylab = "Standardized Residuals",
    xlab = "Happy")
abline(h = 0, lty = 2, lwd = 2)
```

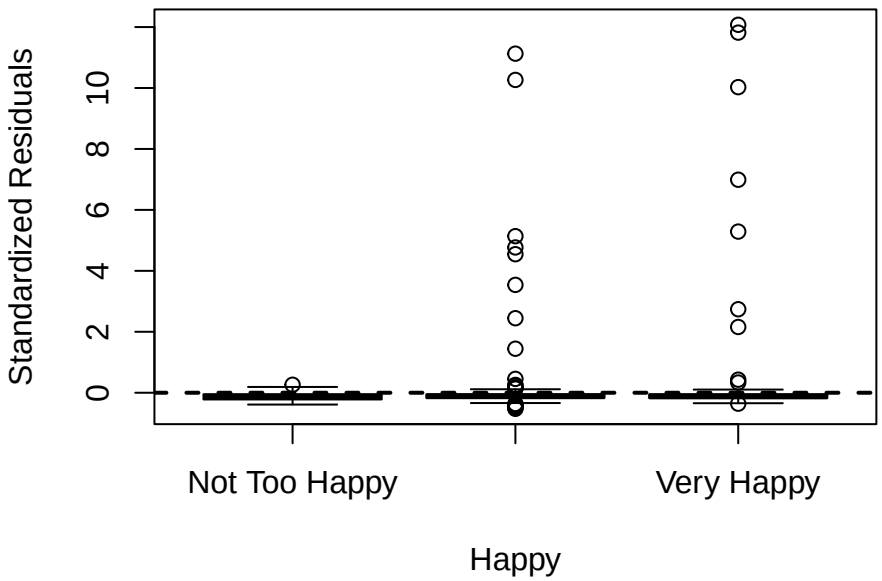

Marital status

```
boxplot(rstandard(full_model) ~ Marital, data = gss, ylab = "Standardized Residuals",
    xlab = "Marital")
abline(h = 0, lty = 2, lwd = 2)
```

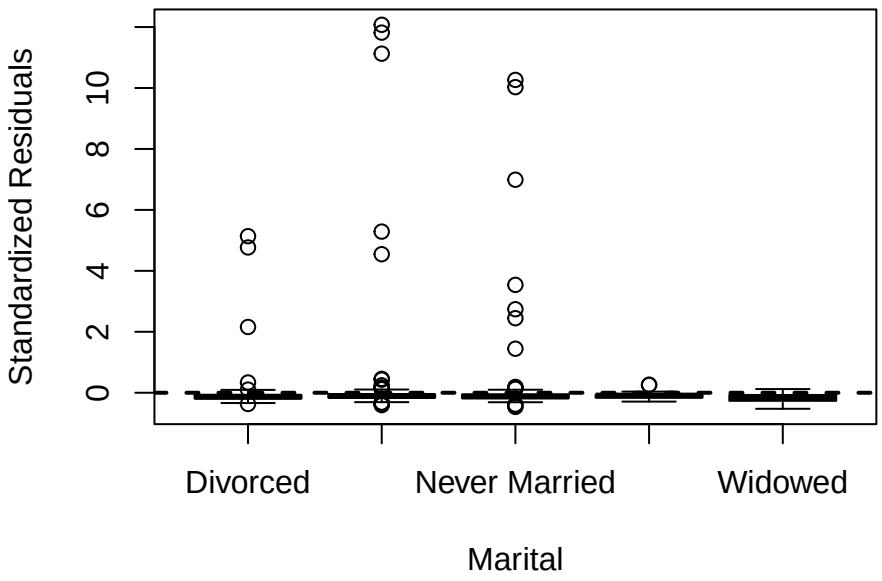

Work status

```
boxplot(rstandard(full_model) ~ Work, data = gss, ylab = "Standardized Residuals",
    xlab = "Work")
abline(h = 0, lty = 2, lwd = 2)
```

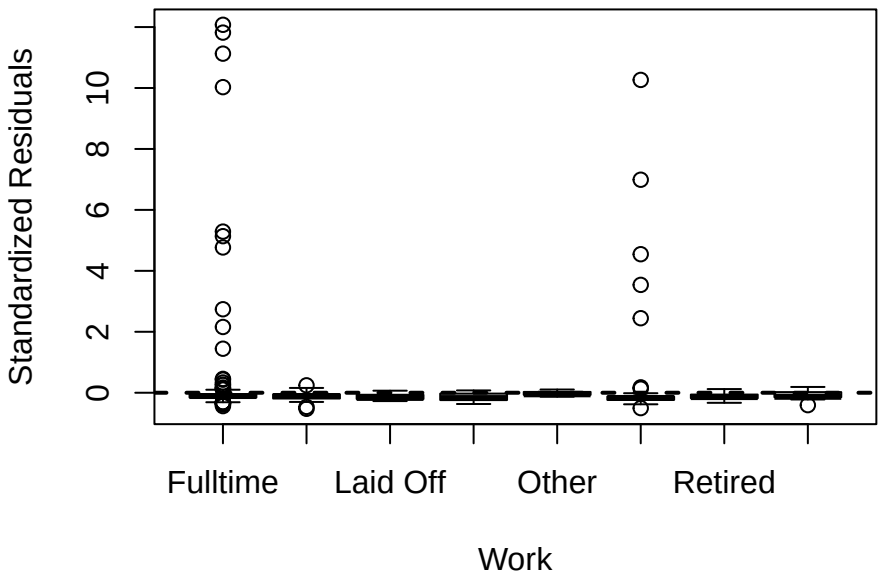

Zodiac sign

```
boxplot(rstandard(full_model) ~ Zodiac, data = gss, ylab = "Standardized Residuals",
    xlab = "Zodiac")
abline(h = 0, lty = 2, lwd = 2)
```

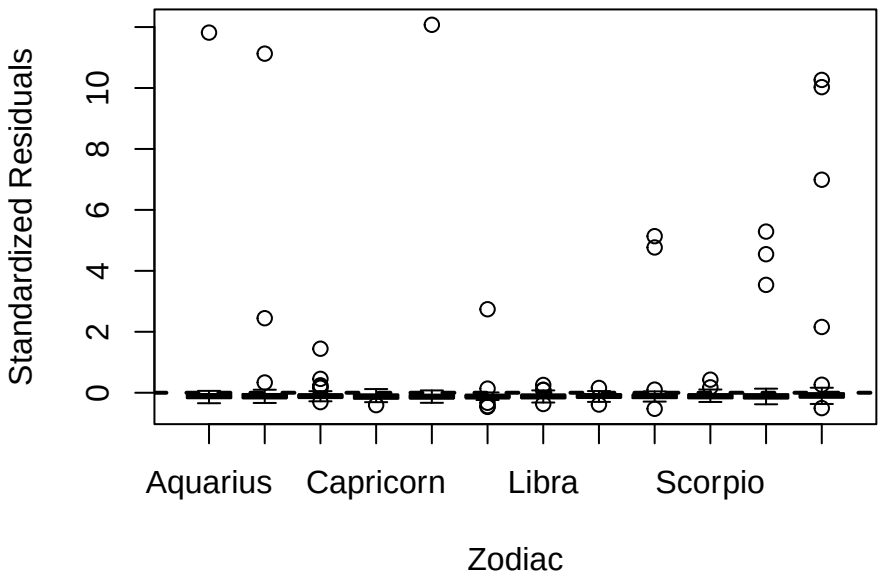

After inspecting the standardized residuals plotted against each qualititative covariate, I again noticed the presence of extreme outliers but no pattern of difference in spread across the levels of the qualitative covariate that would indicate the form of heteroscedasticity that appears when the variance of the random error terms differs across the different levels of some qualitative covariate.

So no, we do not have strong evidence to conclude that the variance of the random error terms is non-constant as a function of any of the quantitative or qualitative covariates.

ii) From the the Breush-Pagan test for heteroscedasticity, our test statistic is 8.955868 and p-value is 0.2558401. Thus, we can not reject the null hypothesis that the full_model is homooscedastic since the p-value is larger than the alpha threshold of 0.05. This means we do not have sufficient evidence to reject the null hypothesis, so we cannot call the data heteroscedastic. This could be done because of the leverage outliers that is overriding the test, causing the test see random noise rather than a linear increase in variance, which leads to a high p-value. That is a possible explanation for this.

```
# Breusch-Pagan Test for heteroscedasticity
Y = gss$Income
residuals = full_model$residuals
Yaux = residuals^2
aux = lm(Yaux ~ Age + Educ + Hours + TVHours + MEduc + FEduc + Sibs,
    data = gss)
BP = n * summary(aux)$r.squared
print(BP)
```

```
## [1] 8.955868
```

```r
pvalue = pchisq(BP, p, lower.tail = FALSE)
print(pvalue)
```

```
## [1] 0.2558401
```

```r
library(lmtest)
bptest(full_model)
```

```
##
##  studentized Breusch-Pagan test
##
## data:  full_model
## BP = 8.9559, df = 7, p-value = 0.2558
```

Step 3:

i) Calculate a suitable heteroscedasticity-consistent estimate for the covariance matrix of the regression coefficients. How does it compare against the corresponding ordinary least squares estimate?

ii) Calculate heteroscedasticity-consistent t-test statistics and 95 percent CIs for the regression coefficients. How do they compare against the results yielded by the ordinary least squares method?

---

i)

```r
# The third heteroscedasticity-consistent estimate is most robust


# Comparing Robust Standard Errors against the OLS Standard Errors
```

ii)

```r
# T-test for the heteroscedasticity-consistent regression
# coefficient


# 95% CI for the heteroscedasticity-consistent regression
# coefficient


# Overall Comparison
```